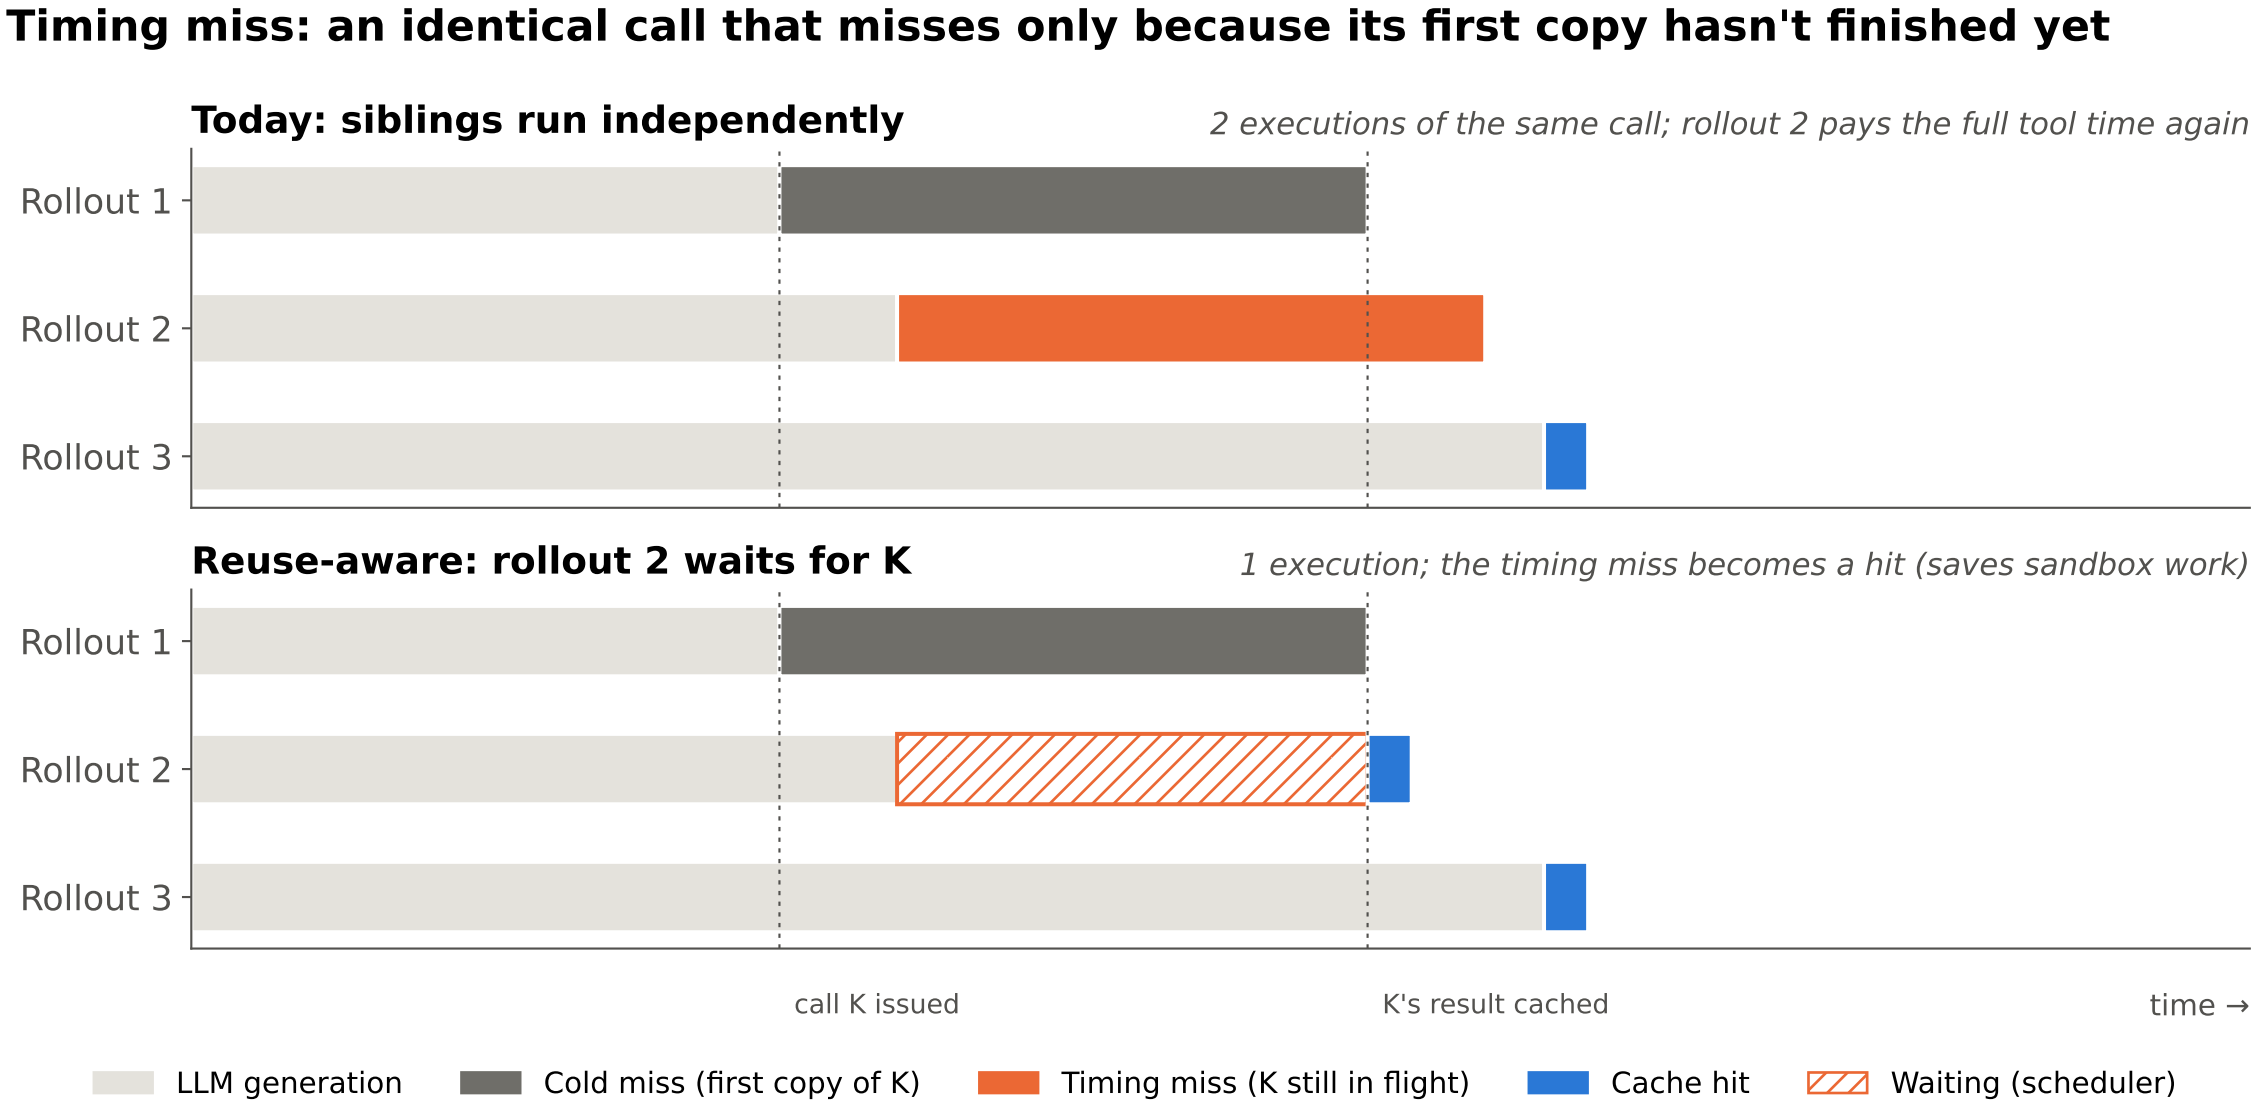
Timing miss: an identical call that misses only because its first copy hasn't finished yet
Today: siblings run independently
2 executions of the same call; rollout 2 pays the full tool time again
Rollout 1
Rollout 2
Rollout 3
Reuse-aware: rollout 2 waits for K
1 execution; the timing miss becomes a hit (saves sandbox work)
Rollout 1
Rollout 2
Rollout 3
call K issued
K's result cached
time →
LLM generation
Cold miss (first copy of K)
Timing miss (K still in flight)
Cache hit
Waiting (scheduler)